Likelihood fmri

Parametrisation

This is special parameterisation of the non-central χ_ν distribution. Let $\{x_i\}$ are iid Gaussians with mean μ and precision τ, then

$$z = \sqrt{\sum_{i=1}^{\nu} \tau x_i^2}$$

is non-central χ-distribution with (integer and fixed by design) $\nu > 0$ degrees of freedom, and non-centrality parameter

$$\rho = \sqrt{\nu \tau \mu^2}.$$

The observation y is $y = z/\sqrt{\tau}$ and we're interested in the underlying true signal

$$\lambda = \rho/\sqrt{\tau} = \sqrt{\nu \mu^2}$$

Link-function

The linkfunction is given as

$$\log(\lambda) = \eta$$

where η is the linear predictor.

Hyperparameters

The hyperparameters are $\theta = (\theta_1, \theta_2)$, where

$$\tau = \exp(\theta_1)$$

is the precision, and

$$\nu = \theta_2$$

For technical reasons, ν is implemented as a hyper-parameter, but is required to be fixed. Hence, the initial value for θ_2 defines the (fixed) value for ν.

The prior is given on θ_1.

Specification

- `family="fmri"` or `family="fmrisurv"`

- Required arguments: y (and optional `scale` for `fmri` to scale τ)

Hyperparameter spesification and default values

doc `fmri distribution (special nc-chi)`

hyper

> **theta1**
>
> > **hyperid** `103101`
> > **name** `precision`
> > **short.name** `prec`
> > **output.name** `Precision for fmri`
> > **output.name.intern** `Log precision for fmri`

       initial 0

       fixed FALSE

       prior loggamma

       param 10 10

       to.theta function(x) log(x)

       from.theta function(x) exp(x)

    theta2

       hyperid 103202

       name dof

       short.name df

       output.name NOT IN USE

       output.name.intern NOT IN USE

       initial 4

       fixed TRUE

       prior normal

       param 0 1

       to.theta function(x) x

       from.theta function(x) x

status experimental

survival FALSE

discrete FALSE

link default log

pdf fmri

doc fmri distribution (special nc-chi)

hyper

    theta1

       hyperid 104101

       name precision

       short.name prec

       output.name Precision for fmrisurv

       output.name.intern Log precision for fmrisurv

       initial 0

       fixed FALSE

       prior loggamma

       param 10 10

       to.theta function(x) log(x)

       from.theta function(x) exp(x)

    theta2

       hyperid 104201

       name dof

        short.name `df`

        output.name `NOT IN USE`

        output.name.intern `NOT IN USE`

        initial `4`

        fixed `TRUE`

        prior `normal`

        param `0 1`

        to.theta `function(x) x`

        from.theta `function(x) x`

status `experimental`

survival `TRUE`

discrete `FALSE`

link `default log`

pdf `fmri`

Example

In the following example we estimate the parameters in a simulated example.

```
n <- 300
x <- rnorm(n, sd = 0.3)
df <- 1
prec <- 3
eta <- 1 + x
lambda <- exp(eta)
y <- sqrt(rchisq(n, df = df, ncp = prec * lambda^2) /prec)

r <- inla(y ~ 1 + x,
          data = data.frame(y, x),
          family = "fmri",
          control.family = list(hyper = list(df = list(initial = df))),
          control.inla = list(cmin = 0,
                              int.strategy = "eb",
                              strategy = "adaptive"),
          verbose = TRUE)
summary(r)

## 'cmin=0' seems to be required only for initial values that can give
## 'crazy' values. We can rerun without this re-starting at the prev fit,
## to validate
r$.args$control.inla$cmin <- -Inf
r$.args$control.inla$int.strategy <- "auto"
rr <- inla.rerun(r)
summary(rr)
```

Notes

Thanks to LS for providing all the details and a robust implementation of this likelihood.