Likelihood fmri

Parametrisation

This is special parameterisation of the non-central χ_ν distribution. Let $\{x_i\}$ are iid Gaussians with mean μ and precision τ, then

$$z = \sqrt{\sum_{i=1}^{\nu} \tau x_i^2}$$

is non-central χ-distribution with (integer and fixed by design) $\nu > 0$ degrees of freedom, and non-centrality parameter

$$\rho = \sqrt{\nu \tau \mu^2}.$$

The observation y is $y = z/\sqrt{\tau}$ and we're interested in the underlying true signal

$$\lambda = \rho/\sqrt{\tau} = \sqrt{\nu \mu^2}$$

Link-function

The linkfunction is given as

$$\log(\lambda) = \eta$$

where η is the linear predictor.

Hyperparameters

The hyperparameters are $\theta = (\theta_1, \theta_2)$, where

$$\tau = \exp(\theta_1)$$

is the precision, and

$$\nu = \theta_2$$

For technical reasons, ν is implemented as a hyper-parameter, but is required to be fixed. Hence, the initial value for θ_2 defines the (fixed) value for ν.

The prior is given on θ_1.

Specification

- family = `fmri` or family = `fmrisurv`

- Required arguments: y (and optional `scale` for `fmri` to scale τ)

Hyperparameter spesification and default values

doc fmri distribution (special nc-chi)

hyper

> **theta1**
>> **hyperid** 103101
>> **name** precision
>> **short.name** prec
>> **initial** 0
>> **fixed** FALSE

prior loggamma

param 10 10

to.theta `function(x) log(x)`

from.theta `function(x) exp(x)`

    theta2

        hyperid 103202

        name dof

        short.name df

        initial 4

        fixed TRUE

        prior normal

        param 0 1

        to.theta `function(x) x`

        from.theta `function(x) x`

status experimental

survival FALSE

discrete FALSE

link default log

pdf fmri

doc fmri distribution (special nc-chi)

hyper

    theta1

        hyperid 104101

        name precision

        short.name prec

        initial 0

        fixed FALSE

        prior loggamma

        param 10 10

        to.theta `function(x) log(x)`

        from.theta `function(x) exp(x)`

    theta2

        hyperid 104201

        name dof

        short.name df

        initial 4

        fixed TRUE

        prior normal

        param 0 1

        to.theta `function(x) x`

from.theta `function(x) x`

status experimental

survival TRUE

discrete FALSE

link default log

pdf fmri

Example

In the following example we estimate the parameters in a simulated example.

```
n <- 300
x <- rnorm(n, sd = 0.3)
df <- 1
prec <- 3
eta <- 1 + x
lambda <- exp(eta)
y <- sqrt(rchisq(n, df = df, ncp = prec * lambda^2) /prec)

r <- inla(y ~ 1 + x,
          data = data.frame(y, x),
          family = "fmri",
          control.family = list(hyper = list(df = list(initial = df))),
          control.inla = list(cmin = 0,
                              int.strategy = "eb",
                              strategy = "adaptive"),
          verbose = TRUE)
summary(r)

## 'cmin=0' seems to be required only for initial values that can give
## 'crazy' values. We can rerun without this re-starting at the prev fit,
## to validate
r$.args$control.inla$cmin <- -Inf
r$.args$control.inla$int.strategy <- "auto"
rr <- inla.rerun(r)
summary(rr)
```

Notes

Thanks to LS for providing all the details and a robust implementation of this likelihood.